

Architecting Control: A Framework for Secure Agent Configuration Management, Approval, and Auditing

Part I: Foundational Principles of Secure Configuration Management

The management of agent configurations represents a critical control plane within any modern enterprise. These agents, whether for security monitoring, performance management, or infrastructure automation, often operate with significant privileges. An unauthorized or erroneous modification can lead to catastrophic consequences, including system-wide outages, data breaches, or the complete circumvention of security controls. Therefore, establishing a robust framework for controlling who can modify these configurations, how changes are approved, and how every action is logged is not merely a technical task but a strategic imperative. This framework must be built upon a set of foundational principles that prioritize security and accountability at every stage. The cornerstone of this entire structure is the Principle of Least Privilege, which informs every subsequent decision regarding access models, approval workflows, and audit requirements.

1.1 The Principle of Least Privilege (PoLP) as the Cornerstone of Access Control

The Principle of Least Privilege (PoLP) is a foundational computer security concept dictating that a user, application, or system should be granted only the bare minimum levels of access—or permissions—necessary to perform its required functions.¹ It is the practice of restricting access rights to the absolute minimum required for the performance of routine, authorized tasks. This principle is not a discretionary guideline but the essential starting point for architecting a defensible security posture for agent configuration management. Its implementation is vital for managing an organization's exposure to cyber threats by fundamentally restricting the potential

damage that can be inflicted by a compromised account, a malicious insider, or a simple human error.¹

The strategic benefits of rigorously applying PoLP are substantial and multi-faceted:

- **Minimizing the Attack Surface:** Every permission granted to a user or system represents a potential avenue for attack. By ensuring that entities are not assigned unnecessary access, PoLP systematically reduces the overall attack surface of the organization.¹ If an attacker compromises an account, the scope of their potential actions is inherently limited to the minimal set of permissions assigned to that account, making it significantly more difficult for them to achieve their objectives.¹
- **Containing Malware and Breaches:** A key tactic for attackers after gaining an initial foothold is lateral movement—moving from the compromised system to other connected devices and systems across the network. PoLP is a powerful countermeasure to this tactic. By limiting an account's privileges, malware or an attacker's actions are often confined to the initial point of entry, preventing the spread of the breach and containing the damage.⁵
- **Mitigating Human Error and Insider Threats:** Not all threats are external. Privileged accounts are a primary target for malicious insiders and third-party vendors seeking to exploit their access.³ Furthermore, accidental misuse of privileges by well-intentioned employees can lead to significant operational disruptions or security gaps. PoLP safeguards against both scenarios by ensuring that no single individual has excessive permissions that could be abused or misused, whether through mistake, malice, or negligence.⁵
- **Meeting Compliance and Zero Trust Requirements:** PoLP is a core component of a Zero Trust security model, which operates under the assumption that no entity, whether inside or outside the network, should be trusted by default.³ Within a Zero Trust framework, PoLP ensures that even after an entity is authenticated, its permissions are strictly limited to its role, and every access request is verified.¹ This alignment makes PoLP a critical element for meeting the requirements of numerous regulatory and industry security standards.

1.2 Implementing PoLP: From Audits to AI-Enhanced Monitoring

The implementation of the Principle of Least Privilege is not a single action but a continuous lifecycle of assessment, enforcement, and monitoring. A successful PoLP

program requires a systematic approach that begins with understanding the current state and evolves into a dynamic, technology-assisted process.

The journey starts with a comprehensive **privilege audit**. This initial step involves discovering and inventorying all privileged accounts across the entire enterprise, including on-premises systems, cloud assets, and DevOps environments.³ The goal is to create a complete picture of the current state of privilege, determining which permissions are assigned to various users and what permissions are actually necessary to perform different tasks.¹ This audit establishes the baseline from which all subsequent PoLP efforts are measured.

Following the audit, the core enforcement tenet is to **default to minimal privileges**. All new user accounts, without exception, must be created with the lowest possible level of access as the default setting.² Privileges should then be granted explicitly and incrementally based on documented business or operational needs. Critically, this process must be symmetrical: when a user's role changes or a project is completed, any privileges that are no longer required must be promptly revoked. This practice is essential to prevent "privilege creep," a common security vulnerability where users accumulate excessive permissions over time, leaving a trail of over-privileged accounts that are prime targets for attackers.²

To further structure access, organizations should implement **segregation of privileges**. This involves creating hard boundaries between different classes of accounts. A common model includes three main types²:

1. **Superuser Accounts:** Often called "admin" accounts, these possess the highest level of privilege and should be reserved for a minimal number of trusted administrators responsible for network monitoring and maintenance.
2. **Least-Privileged User Accounts (LPUs):** These accounts have the bare minimum privileges necessary for routine tasks and should be the standard for the vast majority of employees.
3. **Guest User Accounts:** These have even fewer privileges than LPUs and are granted for limited, temporary access.

Finally, PoLP is not a "set and forget" policy. It demands **continuous monitoring and AI-enhancement**. Organizations must continuously monitor endpoints and audit user accounts to ensure that privilege levels remain appropriate.² Modern approaches leverage artificial intelligence and machine learning (AI/ML) to augment this monitoring. AI/ML systems can establish a baseline of normal access behavior for each user and role. They can then identify anomalous access requests in real-time, such as a user attempting to leverage permissions they have but do not normally use.

Such an event could be an early indicator of a compromised account being exploited by an attacker, allowing for a much faster security response.¹

1.3 Access Control Models: A Comparative Analysis

Implementing the Principle of Least Privilege requires a formal access control model. The two predominant models are Role-Based Access Control (RBAC) and Attribute-Based Access Control (ABAC). These are not mutually exclusive competitors but rather distinct tools, each with its own strengths, weaknesses, and ideal use cases. The choice between them has profound implications for an organization's security posture, administrative overhead, and even the nature of its audit data.

1.3.1 Role-Based Access Control (RBAC): Structure, Scalability, and Limitations

Role-Based Access Control (RBAC) is an access control approach that bases permissions on a user's defined role within an organization.⁶ In this model, permissions—such as the ability to read, write, or execute—are not assigned directly to individual users. Instead, they are aggregated into roles (e.g., "Agent_Configuration_Admin," "Security_Auditor," "System_Operator"), and users are then assigned to these roles.¹ This simplifies administration, as permissions can be managed for a handful of roles rather than for hundreds or thousands of individual users.⁶

The core elements of any RBAC system are ⁶:

- **Administrators:** The users responsible for defining roles and assigning permissions to them.
- **Roles:** Groups of users defined by their job function or responsibilities.
- **Permissions:** The specific actions and access rights granted to each role.

The National Institute of Standards and Technology (NIST) has defined several models of RBAC to suit different organizational needs ⁶:

- **Flat RBAC:** The simplest model, where each user is assigned one or more roles. To gain new privileges, a user must be explicitly assigned an additional role.
- **Hierarchical RBAC:** This model establishes a seniority structure where senior

roles automatically inherit all the permissions of their subordinate roles, in addition to their own unique permissions. For example, a "Senior_Config_Admin" would inherit all permissions of a "Junior_Config_Admin."

- **Constrained RBAC:** This model introduces the concept of segregation of duties, preventing the assignment of conflicting roles to a single user.
- **Symmetric RBAC:** This model mandates that permissions and role assignments are periodically reviewed and adjusted, helping to combat privilege creep.

The primary advantage of RBAC is its straightforwardness and ease of administration, particularly in organizations with stable, well-defined hierarchical structures.⁷ It is relatively easy to implement and automate.⁸ However, its main disadvantage is a phenomenon known as "role explosion." In large or dynamic organizations, the need for fine-grained permissions can lead to the creation of hundreds or even thousands of highly specific roles, which ultimately becomes as difficult to manage as individual permissions.⁷ RBAC struggles in environments where access needs are not cleanly aligned with job titles and change frequently based on projects or tasks.

1.3.2 Attribute-Based Access Control (ABAC): Dynamic, Granular, and Context-Aware Control

Attribute-Based Access Control (ABAC) represents a more dynamic and granular paradigm. Instead of relying on static roles, ABAC grants or denies access by evaluating a set of policies or rules against the attributes of the user, the resource, the desired action, and the environment.¹ This allows for real-time, context-aware access decisions.⁶

An ABAC decision combines several elements⁷:

- **Subject Attributes:** Properties of the user requesting access (e.g., job title, department, security clearance, training certifications).
- **Object/Resource Attributes:** Properties of the resource being accessed (e.g., data classification, owner, creation date).
- **Action Attributes:** The specific action the user is attempting to perform (e.g., read, write, delete, approve).
- **Environmental Attributes:** Contextual factors surrounding the access request (e.g., time of day, geographic location, IP address, device security posture, current threat level).

The key advantage of ABAC is its immense flexibility and power. It enables fine-grained control that can enforce complex business and security rules, such as "Allow marketing managers to edit the 'Campaign_Config' file only from a corporate device within their designated country during normal business hours".⁷ This makes ABAC exceptionally well-suited for complex, distributed organizations with hybrid workforces or temporary project teams. It is generally considered more secure than RBAC because it can prevent access even if an attacker has stolen valid credentials, by evaluating the environmental context of the request and identifying it as anomalous.⁸

However, this power comes at the cost of complexity. Designing and implementing an ABAC system is significantly more time-consuming and requires deep expertise to define the necessary attributes and write the access policies correctly from the outset.⁷ The real-time evaluation of multiple attributes for every access request can also introduce performance latency, especially as the number of users and policies grows.⁸

A critical consideration often overlooked when choosing between these models is the direct and significant impact on the complexity and volume of the resulting audit trail. An RBAC-based access decision is relatively simple to log: "User U, possessing Role R, performed Action A on Object O." The log entry is concise because the context is static and embedded within the role definition. In contrast, an ABAC-based decision is dynamic and highly contextual. The corresponding audit log must capture not only the user and action but the entire set of attributes and environmental conditions that were evaluated by the policy engine to arrive at the "allow" or "deny" decision. This includes the user's location, the device's health, the time of day, and any other relevant attributes. Consequently, an ABAC system generates logs that are far more verbose and complex. This has major downstream implications for the cost of log storage, the complexity of SIEM parsing rules required to make sense of the data, and the level of effort needed for effective post-incident analysis. A leader choosing ABAC for its superior security must simultaneously plan and budget for a more robust and sophisticated logging, storage, and analysis infrastructure to handle the data it produces.

1.4 Choosing the Right Model: A Risk-Based Decision Framework

The choice between RBAC and ABAC is not a matter of one being universally "better" than the other; it is a strategic decision that must be based on an organization's specific context, risk appetite, and technical maturity.

- **Organizational Complexity and Structure:** For smaller organizations or those with a simple, stable, and hierarchical structure, RBAC is often the more efficient and manageable choice. Its simplicity aligns well with clearly defined job functions.⁷ Conversely, large, complex, or highly dynamic organizations—such as those with distributed workforces, a heavy reliance on temporary project teams, or a flat organizational structure—will derive greater benefit from the flexibility and granularity of ABAC.⁷
- **Risk Profile and Compliance Needs:** Industries with a high-risk profile or stringent regulatory requirements, such as finance and healthcare, often find that the granular, context-aware control of ABAC is necessary to meet their security and compliance obligations.⁸ ABAC's ability to enforce rules based on data sensitivity and environmental context provides a stronger security posture for protecting critical assets.
- **The Hybrid Approach:** For many organizations, the optimal solution is not a binary choice but a hybrid approach. RBAC can be used to establish broad, baseline access permissions for the majority of the workforce, providing a manageable foundation. ABAC can then be layered on top for specific high-risk systems, sensitive data categories, or privileged operations. This allows the organization to benefit from RBAC's administrative simplicity for general access while leveraging ABAC's powerful, granular control where it is needed most.

Furthermore, the rise of AI-powered security monitoring creates a feedback loop that favors the adoption of richer access control models. AI and machine learning systems designed to detect anomalous access patterns thrive on rich, contextual data.¹ An ABAC system, by its very nature, provides this detailed context with every access request—capturing user, resource, action, and environmental attributes. An RBAC system provides far less contextual information. Therefore, an organization that implements ABAC is better positioned to leverage the full potential of next-generation, AI-driven security monitoring. The AI can build more sophisticated and accurate models of "normal" behavior, leading to fewer false positives and a greater ability to detect subtle threats. This creates a virtuous cycle: ABAC enables better AI analysis, which in turn can provide data-driven recommendations for refining and strengthening the ABAC policies, leading to a more adaptive and intelligent security posture.

The following table provides a strategic comparison to guide the decision-making

process.

Table 1: RBAC vs. ABAC - A Strategic Comparison

Criterion	Role-Based Access Control (RBAC)	Attribute-Based Access Control (ABAC)
Granularity	Coarse-grained; permissions tied to static roles.	Fine-grained; permissions based on dynamic attributes of user, resource, action, and environment. ⁸
Scalability	Scales well in stable organizations but can suffer from "role explosion" in complex ones. ⁷	Highly scalable for complex and dynamic environments, but policy management can become complex. ⁸
Implementation	Relatively simple and straightforward to configure and deploy. ⁷	Complex and time-consuming; requires significant upfront effort to define attributes and policies. ⁷
Administrative Overhead	Low for managing roles, but can become high if role explosion occurs.	High during initial policy creation, but can simplify management of dynamic access needs over time.
Audit Trail Complexity	Generates simpler, less verbose logs. "User in Role performed Action."	Generates complex, verbose logs capturing all evaluated attributes and context. Requires more advanced analysis tools. ¹
Ideal Use Case	Small to medium-sized organizations with stable, hierarchical structures and well-defined job functions. ⁷	Large, complex, or distributed organizations; high-security environments; industries with strict compliance needs. ⁷
Synergy with AI Monitoring	Provides limited context for behavioral analysis.	Provides rich, contextual data that significantly enhances the effectiveness of AI/ML-based anomaly detection. ¹

Part II: Designing Robust Approval and Change Management Workflows

Once the principles of access control have established *who* can potentially modify agent configurations, the next critical layer of defense is the procedural framework governing *how* those modifications are made. For changes that can impact system stability, security posture, or risk exposure, a simple grant of privilege is insufficient. A formal, documented, and enforced change management and approval workflow is essential. This part addresses the user's direct query about multi-person approval processes for critical changes, such as increasing risk limits, by detailing the principles of Segregation of Duties and the practical implementation of the "maker-checker" workflow.

2.1 The Imperative for Segregation of Duties (SoD) in Critical Changes

Segregation of Duties (SoD), sometimes referred to as Separation of Duties, is a foundational concept in internal controls designed to prevent fraud, abuse, and errors. The principle states that a critical task should be divided among multiple individuals to ensure that no single person has end-to-end control over the process.¹⁰ By separating responsibilities, an organization creates a system of checks and balances that makes it impossible for one person to commit a malicious act or a catastrophic error without being detected by another.¹⁰

When applied to the management of agent configurations, SoD is not an abstract ideal but a mandatory control, especially for high-risk changes. For a critical modification, such as increasing a financial risk limit, changing a security policy, or altering a data retention setting, the principle of SoD dictates that the person who requests or proposes the change cannot be the same person who approves and implements it. This separation is a cornerstone of compliance with numerous regulatory frameworks, most notably the Sarbanes-Oxley Act (SOX), which places stringent requirements on internal controls over financial reporting systems.¹¹ A failure to implement SoD for critical system changes is a significant control deficiency that

would be flagged by any competent auditor.

2.2 The Maker-Checker (Four-Eyes) Principle: A Procedural Deep Dive

The "maker-checker" principle, also known as the "four-eyes" principle or two-man rule, is the direct and practical implementation of Segregation of Duties in a transactional workflow.¹³ It is a central principle of authorization, particularly in financial organizations, and its logic is directly applicable to the control of critical agent configurations.¹⁵ The core concept is simple yet powerful: any sensitive activity must be initiated by one person (the "maker") and then independently verified and authorized by a second person (the "checker") before it can be completed.¹⁴

The benefits of a maker-checker workflow are clear and compelling:

- **Prevents Internal Fraud:** By requiring the action of two separate individuals, the system makes it significantly harder for a malicious insider to abuse their privileges. Committing fraud would require collusion between the maker and the checker, raising the difficulty and risk of detection.¹⁵
- **Reduces Human Error:** A second, independent review by the checker provides a crucial opportunity to catch configuration errors, typos, or logical flaws that the maker might have overlooked. This "second pair of eyes" is a simple but highly effective method for improving quality control and preventing accidental outages or security vulnerabilities.¹⁵
- **Increases Transparency and Accountability:** The maker-checker process inherently creates a clear, documented workflow. Every step of the process is recorded, establishing an unambiguous chain of accountability for every critical change.¹³

A well-designed maker-checker workflow for agent configuration changes should follow a defined sequence of steps:

1. **Request Initiation (Maker):** A user with designated "maker" privileges submits a formal request for a configuration change. This request must include not only the proposed technical change but also a clear business justification, an assessment of potential impact, and a reference to a corresponding change ticket or project ID.
2. **Verification and Approval (Checker):** The change request is automatically routed to one or more users with "checker" privileges. Crucially, the system must

enforce that the checker cannot be the same individual as the maker. The checker is responsible for reviewing the entire request—the technical accuracy of the change, the validity of the justification, and the assessment of risk.

3. **Execution:** Only after the checker has formally approved the request is the change cleared for implementation. Depending on the system's design, this implementation could be performed manually by a third role (e.g., a system operator) or, in a more advanced setup, triggered automatically by the approval system.
4. **Documentation and Auditing:** The entire lifecycle of the request—the initial submission, the maker's identity, the checker's identity, all associated timestamps, comments, and the final outcome (approved or rejected)—must be immutably recorded in the system's audit trail.¹⁶

2.3 Automating Approval Workflows: Balancing Efficiency and Control

While the maker-checker principle can be implemented manually through emails, chat messages, or paper forms, such methods are notoriously inefficient, prone to error, and lack the robust, centralized audit trail required by regulators.¹⁷ Manual processes create bottlenecks when approvers are unavailable, documents can be lost or misfiled, and it becomes difficult to track the status of any given request.

The modern solution is to leverage workflow automation software to design, enforce, and document approval processes.¹⁷ Automating the approval workflow provides several key advantages:

- **Standardization and Consistency:** An automated system ensures that every change request follows the same predefined, standardized steps, eliminating inconsistencies in how approvals are handled across different teams or departments.¹⁷
- **Efficiency and Speed:** Automation streamlines the process by automatically routing requests to the correct approvers, sending reminders for pending tasks, and providing notifications upon completion. This significantly reduces delays and bottlenecks.¹⁷
- **Enhanced Visibility and Control:** Centralized workflow systems provide real-time dashboards where managers and auditors can track the status of all change requests, identify bottlenecks, and review historical approval data.¹⁷
- **Robust Audit Trails:** Automated systems generate comprehensive, consistent,

and easily accessible audit trails for every step of the workflow, simplifying compliance and forensic analysis.¹⁸

When designing an automated workflow, organizations should implement risk-based routing. Not all changes carry the same level of risk. A minor, non-impactful change might be routed through a simplified workflow, perhaps with automated checks and a single approver. A critical change, like modifying a core security policy or increasing a financial transaction limit, should automatically trigger a more stringent, multi-level approval workflow involving multiple checkers from different departments (e.g., IT, Security, and Business). For very low-risk, highly standardized changes, it may even be possible to implement rules for automatic approval based on predefined criteria, further maximizing efficiency without sacrificing control.¹⁷

The adoption of modern development practices like Infrastructure as Code (IaC) and GitOps provides a powerful avenue for embedding these approval workflows directly into the technology toolchain. In this paradigm, the "maker" is the developer who writes the configuration code and submits it as a pull request (PR) or merge request (MR) in a Git repository.¹⁹ The "checker" is the peer reviewer or team lead who must review and formally approve the PR/MR. The Git platform's native features, such as branch protection rules that mandate at least one approval before code can be merged into the main branch, become the technical enforcement mechanism for the maker-checker principle.¹⁹ This transforms the approval workflow from a separate, manual process into an integrated, automated, and auditable part of the development lifecycle itself. The CISO's role thus shifts from policing a manual process to ensuring that the development platforms are configured with the correct automated guardrails.

2.4 Managing Emergency Changes: The "Break-Glass" Procedure

While structured approval workflows are essential for routine changes, there will inevitably be emergency situations—such as a critical production outage or an active security incident—where waiting for a standard multi-person approval process is not feasible. For these scenarios, a formal, documented, and highly controlled "break-glass" procedure is required. This procedure allows for temporary, emergency access to make urgent changes, but it must be governed by strict compensating controls.

A robust break-glass procedure should include the following elements:

1. **Formal Request and Justification:** The user requiring emergency access must initiate a formal "break-glass" request, explicitly citing the reason for the emergency (e.g., an incident ticket number). This creates an initial record of intent.
2. **Temporary and Specific Privilege Elevation:** The system should grant the user elevated privileges that are strictly temporary and scoped to the minimum necessary to resolve the specific incident. Broad, permanent administrator access should not be granted.
3. **Intensive, Verbose Logging:** All actions performed by the user during the break-glass session must be logged at the highest possible level of detail. This logging should be tamper-evident and reviewed as a high-priority item.
4. **Mandatory Post-Facto Review:** This is the most critical compensating control. After the emergency has been resolved, there must be a mandatory, documented review of all activities that occurred during the break-glass session. This review should be conducted by management and the security team to ensure that the emergency access was used appropriately and only for its stated purpose.²¹
5. **Automatic Privilege Revocation:** The elevated privileges must be revoked automatically after a short, predefined time limit (e.g., a few hours) to ensure that the emergency access does not inadvertently become permanent.

There is an inherent and unavoidable tension between the velocity demanded by modern DevOps practices and the deliberate friction introduced by formal approval workflows. DevOps and CI/CD pipelines are designed to "fail fast and fix fast," prioritizing speed and automation.²² Approval workflows, by their very nature, introduce a pause for human review and judgment, which is essential for risk management.¹⁷ Attempting to eliminate this friction entirely in the name of speed would negate the purpose of the control. The goal, therefore, is not to eliminate approvals but to make the process as efficient and frictionless as possible. This involves deeply integrating the approval mechanism into the developer's native workflow, automating the routing and notification of requests, and providing reviewers with all the necessary context—such as automated security scan results and impact analyses—directly within their review interface. The objective is to achieve "secure velocity," where controls are embedded in a way that minimizes delay and cognitive load while maintaining the integrity and purpose of the approval gate.

Part III: The Immutable Record: Architecting Comprehensive and Tamper-Evident Audit Trails

The third pillar of a secure configuration management framework is the audit trail. After establishing who can make changes (access control) and how those changes are approved (change management), it is imperative to create a complete, accurate, and trustworthy record of every single action taken. A comprehensive audit trail is not just a compliance checkbox; it is an essential tool for post-incident forensic analysis, operational troubleshooting, and demonstrating accountability to auditors and regulators. This section details the required components of an effective log entry and explores the technological progression from basic log protection to advanced cryptographic techniques that can render an audit trail truly tamper-evident and immutable.

3.1 Anatomy of an Effective Audit Log Entry

An audit trail is far more than a simple list of events; it is a detailed, chronological, and context-rich record of all activities within a system.²³ For a configuration change, which can have significant security and operational implications, each log entry must be treated as a piece of critical evidence. To be effective for both auditing and forensic analysis, every log entry related to a configuration change must capture a specific set of data elements.

A modern, automated, and GitOps-driven environment complicates the notion of a single, monolithic log file. The evidence for a single configuration change is often distributed across multiple systems. The initial proposal might be in a project management tool like Jira or ServiceNow. The code change itself is captured as a commit in a Git repository. The approval is recorded in a pull/merge request. The automated deployment process generates logs in a CI/CD platform like Jenkins or GitLab CI. Finally, the effect of the change is recorded in the target application's or system's own audit log. Therefore, a "comprehensive audit trail" is less about creating one perfect log file and more about ensuring the ability to **correlate** these disparate sources of evidence. The final system-level log entry must contain unique identifiers—such as a Git commit hash, a pull request ID, and a change ticket number—that allow an auditor or investigator to trace the entire lifecycle of the change back through the approval and development process.

The following table outlines the essential fields that should be present in the final,

system-level audit log entry for any configuration change.

Table 2: Essential Elements of a Configuration Change Audit Log Entry

Field Name	Description	Example	Source(s)
EventID	A unique, system-generated identifier for the log event itself, ensuring each entry can be uniquely referenced.	evt-a1b2c3d4-e5f6-7890	23
Timestamp	A high-precision, synchronized UTC timestamp indicating when the event occurred, ideally to the millisecond, from a reliable time source (NTP).	2023-10-27T10:00:05.123Z	23
UserID	The unique, non-repudiable identifier of the user or service account that performed the action.	svc_acct_deploy or john.doe	23
SourceIP	The IP address or hostname of the system from which the change was initiated.	10.1.2.3 or build-agent-05.prod.local	23
Action	A clear, standardized verb describing the action taken (e.g., CONFIG_MODIFY, ACCESS_GRANTED, LOGIN_FAILURE).	CONFIG_MODIFY	23

Object	The specific configuration parameter, file, or resource that was the target of the action.	agent.risk.limit	23
OldValue	The value of the configuration parameter <i>before</i> the change was applied. This is critical for rollback and analysis.	100000	23
NewValue	The value of the configuration parameter <i>after</i> the change was applied.	150000	23
Status	The outcome of the action (e.g., SUCCESS, FAILURE). Logging failed attempts is crucial for detecting malicious activity.	SUCCESS	23
ApprovalChainID	The identifier of the approval record in the change management system (e.g., pull request number, workflow ID).	pr-452	21
CorrelationID	An identifier linking this event to a business-level request or incident (e.g., Jira ticket number, ServiceNow Change Request ID).	JIRA-SEC-101	21

3.2 Ensuring Log Integrity: From Basic Protection to Cryptographic Assurance

Capturing comprehensive log data is only the first step. The value of an audit trail is entirely dependent on its integrity. If logs can be modified or deleted without detection, they are worthless as a source of truth. Protecting log integrity is a layered process, progressing from basic operational controls to sophisticated cryptographic guarantees.

Level 1: Foundational Controls (Secure Storage and Access)

The most basic layer of protection involves treating the logs themselves as sensitive assets.

- **Centralized and Remote Logging:** Logs must not be stored on the same systems they are monitoring. An attacker who compromises a system could easily delete local logs to cover their tracks.²⁴ Instead, logs should be immediately transmitted to a secure, centralized log management system or SIEM.²⁶
- **Strict Access Control:** Access to the raw log files and the log management system must be governed by the Principle of Least Privilege. Most users, including system administrators, should only have read-only access. The ability to modify or delete logs should be restricted to an extremely small, highly monitored group of specialized log administrators, if it is allowed at all.²⁴

Level 2: Tamper-Evidence via Cryptographic Chaining

This level moves beyond simple access control to provide active proof of integrity.

- **Hash Chains:** In this model, each new log entry is cryptographically hashed. The input to the hash function for a new entry includes not only the entry's own data but also the hash of the immediately preceding entry. This creates a linked list or "chain" of hashes. If an attacker attempts to alter any historical log entry, its hash will change. This change will cause the hash of the next entry to be incorrect, and this error will cascade through the entire rest of the chain. By periodically checking the hash of the latest entry, an auditor can verify the integrity of the entire chain up to that point.²⁷
- **Digital Signatures:** To provide proof of origin (non-repudiation) and integrity, log entries or batches of entries can be digitally signed using the private key of the system that generated them. Anyone with the corresponding public key can then verify that the log came from the correct source and has not been altered since it was signed.²⁵

Level 3: Advanced Integrity with Merkle Trees

For very large-scale logging systems, hash chains can become inefficient. Merkle trees offer a more advanced and efficient solution for proving integrity.

- **Mechanism:** In this data structure, individual log entries are the "leaf" nodes of a binary tree. Each leaf node is a hash of its log data. Pairs of leaf nodes are then

hashed together to create a parent node, and this process continues up the tree until a single "root hash" is produced. This root hash is a unique cryptographic digest representing the entire set of log entries.²⁵

- **Advantages over Hash Chains:**

- **Efficient Proof of Inclusion:** To prove that a specific log entry exists within a dataset of millions or billions of entries, one does not need to re-hash the entire chain. Instead, a small "Merkle proof" can be provided. This proof consists of the sibling hashes along the path from the leaf to the root. An auditor can use this small proof (whose size is logarithmic, $O(\log N)$, relative to the number of logs) to independently recalculate the root hash and verify the entry's inclusion.²⁷ This is exponentially more efficient than the linear, $O(N)$, effort required for a hash chain.
- **Proof of Consistency:** Merkle trees can also be used to generate efficient proofs that a newer version of the log is a valid continuation of an older version—that is, it only appends new entries without modifying or deleting any old ones. This is crucial for verifying the log's history over time.²⁸

3.3 Achieving True Immutability: Advanced Tamper-Proofing Strategies

The ultimate goal of an audit trail is immutability—the guarantee that once a record is written, it can never be altered or deleted without detection. This requires addressing the most challenging threat model: one where the entire internal infrastructure, including the logging system itself, is compromised. Achieving this level of assurance requires looking beyond internal controls to external, trusted systems.

The Role of Trusted Third-Party Log Storage:

A fundamental shift in the security model occurs when logs are sent to a trusted, external third-party service that is architecturally designed for write-once, read-only storage.²⁵ By removing the logs from the organization's direct control, this strategy mitigates two key risks simultaneously. First, an external attacker who breaches the primary infrastructure cannot pivot to the logging system to tamper with the evidence of their intrusion. Second, it prevents internal actors within the organization from modifying logs to cover up errors, fraud, or compliance failures.²⁵

Client-Side Cryptographic Signing for Non-Repudiation:

To protect against the risk of logs being altered in transit or by the third-party service itself, each log record should be cryptographically signed on the client side, before it is transmitted. This is done using a private key that is known only to the originating system and is never shared. This client-side signature provides irrefutable proof of authorship and ensures that

the record stored by the third party is the exact, unaltered record that was sent.²⁵

Leveraging Public Ledgers (Blockchain) for Verifiable Continuity:

The final piece of the puzzle is to ensure that the third-party logging service itself is behaving honestly and has not lost or altered data. This can be achieved by anchoring the log's integrity in a public, decentralized, and immutable ledger, such as a blockchain.²⁵ The process works as follows:

1. The third-party service maintains the logs in a Merkle tree.
2. Periodically (e.g., every hour), the service takes the latest Merkle root hash.
3. It then publishes this root hash to a public blockchain in a transaction.

This creates a publicly verifiable, time-stamped "proof of continuity." Any independent auditor can retrieve the sequence of root hashes from the blockchain and use them to verify the integrity and completeness of the logs held by the third-party service. Because the blockchain is immutable, neither the organization nor the logging service can go back and change this history.²⁵

While these advanced techniques provide powerful security guarantees, they also introduce new operational risks and dependencies. An organization becomes reliant on the availability and performance of the third-party logging service and the public blockchain network. A robust architecture must therefore include strategies for resilience, such as secure local caching of logs during a service outage, with a mechanism to securely transmit the cached logs once the service is restored. The pursuit of immutability requires careful consideration of the trade-offs between security assurance and operational autonomy.

Part IV: Integrating Controls within Established Governance Frameworks

Implementing robust technical and procedural controls for access, approval, and auditing is essential. However, for these controls to be effective at an enterprise level, they must be aligned with and mapped to established IT governance and cybersecurity frameworks. This alignment serves two critical purposes. First, it ensures that the implemented controls are comprehensive and address risks in a structured, recognized manner. Second, it provides a common language for communicating the value and necessity of these controls to senior leadership, risk managers, and external auditors, demonstrating compliance with industry standards and regulatory requirements.

The leading frameworks—NIST Cybersecurity Framework (CSF), COBIT, and ISO/IEC 27001—are not competing standards but rather different lenses through which to view the same set of underlying security challenges. The NIST CSF provides a framework for managing **risk**. COBIT provides a framework for **governance and management**, defining roles and processes. ISO 27001 provides a standard for an **information security management system (ISMS)**, offering a specific checklist of auditable controls. A well-designed control system for agent configuration management will inherently satisfy the principles of all three. This "Rosetta Stone" approach allows a CISO to efficiently demonstrate a mature, integrated compliance posture to diverse stakeholders.

Table 3: Framework Alignment Matrix (NIST, COBIT, ISO 27001)

Control Area	NIST CSF 2.0 Reference	COBIT 2019 Objective	ISO/IEC 27001:2022 Control
Access Control Policy	GV.PO; PR.AC	APO01; DSS05	A.5.15 Access control
Role Definition & Assignment	PR.AC	DSS05	A.5.16 Identity management
Privilege Management (PoLP)	PR.AC	DSS05	A.5.18 Access rights; A.8.2 Privileged access rights
Change Approval Workflow	GV.PO; PR.PT	BAI06 Manage IT Changes	A.8.32 Change management
Emergency Change Process	PR.PT; RS.MA	BAI06	A.8.32 Change management
Log Generation & Collection	PR.PT; DE.CM	MEA03	A.8.15 Logging
Log Protection & Integrity	PR.DS; PR.PT	DSS05	A.8.16 Monitoring activities

4.1 NIST Cybersecurity Framework (CSF) 2.0

The NIST Cybersecurity Framework provides a voluntary, risk-based approach for organizations to manage and reduce cybersecurity risk.²⁹ CSF 2.0 expands the framework's scope and introduces a new top-level "Govern" function, which underscores the strategic importance of establishing and maintaining the very controls discussed in this report.³¹ The controls for agent configuration management map across multiple CSF functions, demonstrating their central role in a holistic cybersecurity program.

- **Govern (GV):** This new function is the starting point. It encompasses the organization's overall cybersecurity risk management strategy, expectations, and policies.³¹ The decisions to adopt PoLP, implement maker-checker workflows, and mandate tamper-evident logging are all governance-level decisions that set the direction for the entire program.
- **Protect (PR):** This is the primary home for the implementation of these controls. The "Protect" function supports the ability to limit or contain the impact of a potential cybersecurity event.³⁰ Several categories within this function are directly relevant:
 - **PR.AC - Identity Management, Authentication and Access Control:** This category is the direct mapping for the implementation of PoLP, RBAC, and ABAC. Its goal is to ensure that access to assets is limited to authorized users, processes, and devices, consistent with assessed risk.⁹
 - **PR.DS - Data Security:** This category covers the protection of data confidentiality, integrity, and availability. It directly applies to the protection of the audit logs themselves through measures like encryption, access control, and the cryptographic integrity mechanisms discussed in Part III.⁹
 - **PR.PT - Protective Technology:** This category includes the management of technical security solutions. Critically, it lists "audit/log records" as a specific subcategory, reinforcing the need to protect the logging infrastructure itself.³¹ The change management and approval workflows are also part of the protective technology and processes.
- **Detect (DE):** This function enables the timely discovery of cybersecurity events. The comprehensive audit logs generated by the system are a primary input for detection mechanisms. Monitoring these logs for unauthorized change attempts, failed access, or other anomalies is a core detection activity.³¹
- **Respond (RS) & Recover (RC):** In the event of a security incident, the immutable and comprehensive audit trail becomes an indispensable asset. It is used during the response phase for forensic investigation to understand the attacker's actions and during the recovery phase to ensure systems are restored to a known-good

state.³¹

4.2 COBIT 2019

COBIT is an IT governance and management framework that helps organizations align their IT processes with business objectives.³³ A key feature of COBIT is its clear distinction between

Governance (the responsibility of the board and executive management to Evaluate, Direct, and Monitor) and **Management** (the responsibility of operational teams to Plan, Build, Run, and Monitor).³³ The controls for agent configuration span both domains.

- **Governance (Evaluate, Direct and Monitor - EDM):** The board of directors and senior leadership are responsible for directing that a robust internal control system for information and technology is established. This includes setting the risk appetite and policies that mandate strong access control and change management for critical systems like agent configurations.³⁵
- **Management (Align, Plan, Organize - APO; Build, Acquire, Implement - BAI; Deliver, Service, Support - DSS; Monitor, Evaluate, Assess - MEA):** The day-to-day implementation of the controls falls under the management domains.
 - **BAI06 - Manage IT Changes:** This process is the direct COBIT equivalent of the change management workflows detailed in Part II. COBIT provides specific control practices that align perfectly with the required procedures. For example, A16.1 Change standards and procedures calls for establishing a framework that defines roles, responsibilities, authorization, and approval for all changes. A16.2 Change request management details the process for formally requesting and prioritizing changes. These practices directly support the implementation of maker-checker workflows and formal emergency change ("break-glass") procedures.²¹
 - **DSS05 - Manage Security Services:** This domain includes the processes for protecting information systems. This encompasses managing user identity and logical access, which aligns with the PoLP and RBAC/ABAC implementations.
 - **DSS06 - Manage Business Process Controls:** This process involves managing user access rights and ensuring proper segregation of duties, directly mapping to the SoD principles that underpin the maker-checker

workflow.

- **MEA02 - Monitor, Evaluate and Assess the System of Internal Control:** This domain covers the continuous monitoring of controls. The regular review of audit logs to ensure compliance with change management policies and to detect unauthorized activity is a key activity within this process.³⁵

4.3 ISO/IEC 27001:2022

ISO/IEC 27001 is the leading international standard for an Information Security Management System (ISMS). It provides a systematic, risk-based approach to managing an organization's sensitive information.³⁶ The standard's Annex A provides a comprehensive catalogue of 93 potential information security controls, from which an organization selects based on its risk assessment.³⁸ The controls for agent configuration management are explicitly covered in several Annex A controls.

- **A.5 Organizational Controls:** This theme covers the foundational policies and processes for information security.
 - **A.5.15 Access control:** This control requires the organization to establish, document, and review an access control policy based on business and security requirements. This is the top-level policy that mandates PoLP.³⁸
 - **A.5.16 Identity management:** This control requires the management of the full lifecycle of identities, ensuring that every user has a unique identifier to enable accountability.³⁸
 - **A.5.18 Access rights:** This control focuses on the processes for provisioning, reviewing, modifying, and revoking access rights. It mandates that access should be granted based on the principle of least privilege ("need to know" and "need to use").¹⁰
- **A.8 Technological Controls:** This theme covers the technical implementation of security measures.
 - **A.8.2 Privileged access rights:** This control places specific requirements on the management of privileged accounts, mandating that their allocation and use be restricted and monitored.
 - **A.8.5 Secure authentication:** This control requires the use of strong authentication mechanisms to verify the identity of users before granting access.
 - **A.8.15 Logging:** This control mandates that logs recording user activities, exceptions, faults, and information security events shall be produced,

protected, and regularly reviewed. This directly maps to the requirements for a comprehensive audit trail.

- **A.8.32 Change management:** This control is the direct ISO 27001 equivalent for the approval workflows discussed in Part II. It requires that changes to information processing facilities and systems be controlled through a formal change management process. This includes assessing the security impact of changes, obtaining formal approval before implementation, and maintaining a record of all changes.⁴¹

By mapping the implemented controls to these specific requirements from NIST, COBIT, and ISO 27001, an organization can build a coherent, integrated, and defensible security program that stands up to the scrutiny of any audit.

Part V: Modern Implementation and Regulatory Considerations

The foundational principles and governance frameworks provide the "what" and "why" of secure configuration management. This final part addresses the "how" in the context of modern technology stacks and the specific pressures of key regulatory mandates. The paradigm shift towards managing infrastructure as code fundamentally transforms how these controls are implemented, moving them from manual processes to automated guardrails embedded within the development lifecycle. This modern approach not only increases efficiency but also provides a stronger, more auditable control posture that is well-aligned with the stringent requirements of regulations like SOX, HIPAA, and FINRA.

5.1 Configuration as Code: Applying GitOps Principles to Change Management

In a modern cloud-native or automated environment, system and agent configurations are no longer modified manually through a graphical user interface. Instead, they are defined declaratively in text files (e.g., using YAML, HCL) and managed as code, a practice known as Infrastructure as Code (IaC).⁴² This approach brings the power and discipline of software development practices to infrastructure management.

The GitOps methodology takes this a step further by establishing a Git repository as

the single source of truth for the entire desired state of the system.¹⁹ The entire change management lifecycle is managed through a standard Git workflow, which inherently provides a robust framework for control and auditing:

- **Change Request as a Pull Request:** When a developer or operator needs to make a configuration change, they do not log into a server. Instead, they create a new branch in the Git repository, modify the configuration code, and then open a Pull Request (PR) or Merge Request (MR). This PR/MR is the formal, documented request for change, containing the exact code modifications and a description of the change's purpose.¹⁹
- **Approval as Code Review:** The PR/MR triggers a mandatory peer review process. This is the direct, technical implementation of the maker-checker principle. A designated approver (the "checker") reviews the proposed code changes, and only after they provide their formal approval within the Git platform can the change be merged into the main branch. Branch protection rules can be configured to enforce this, for example, by requiring at least one approval from a designated code owner.¹⁹
- **Audit Trail as Git History:** The Git repository itself becomes the primary, cryptographically-secured audit trail. Every commit, every comment in a PR, every approval, and every merge is recorded with a timestamp and attributed to a specific user. This immutable history provides an unparalleled level of transparency and accountability, capturing the who, what, when, and why of every single configuration change automatically.¹⁹

5.2 CI/CD Pipeline Security: Embedding Approval Gates and Scans

Once a change is approved and merged in Git, a Continuous Integration/Continuous Delivery (CI/CD) pipeline automates the process of testing and deploying the new configuration to the production environment.²² This automation provides speed and consistency, but it must be fortified with automated security gates to prevent insecure or erroneous configurations from being deployed. The security team's role shifts from being a manual gatekeeper to being the builder of these automated guardrails.

A secure CI/CD pipeline for configuration management should include several automated security checks:

- **Static Analysis Security Testing (SAST):** Before a PR can even be approved, the pipeline should automatically trigger a scan of the IaC code. SAST tools can

detect common misconfigurations, policy violations, and, critically, hard-coded secrets like API keys or passwords, which are a major security risk.²⁰

- **Policy as Code (PaC):** Tools like Open Policy Agent (OPA) allow security and compliance policies to be written as code. These policies can be automatically enforced by the CI/CD pipeline. For example, a policy could prevent any configuration change that attempts to open a firewall port to the public internet or disable required logging.
- **Automated Build Failure:** The pipeline must be configured to **fail the build** if any of these automated security checks do not pass or if the required approvals have not been granted. This is a critical control that programmatically prevents an insecure or unapproved change from ever reaching the production environment, effectively shifting security "left" into the earliest stages of the development process.²²

This shift to a GitOps model represents a significant cultural and operational change. The security team is no longer a downstream checkpoint that can be perceived as a bottleneck. Instead, they become upstream enablers and architects of the secure "paved road." Their primary function is to provide developers with the tools, policies, and automated guardrails that allow them to innovate quickly but safely. This requires a new skillset for security professionals, blending deep security knowledge with software development and automation expertise to build and maintain these critical pipeline controls.

5.3 Regulatory Crosswalk: Mapping Controls to Specific Mandates

The robust, automated, and highly auditable nature of a modern, GitOps-driven configuration management framework aligns exceptionally well with the requirements of major regulatory mandates. While these regulations were often written before such technologies were widespread, their core principles of integrity, accountability, and audibility are better met by these modern approaches than by traditional manual controls.

- **Sarbanes-Oxley (SOX):** SOX compliance hinges on the effectiveness of IT General Controls (ITGCs) that protect the integrity of systems involved in financial reporting.¹² The key ITGC domains of **change management** and **access controls** are directly and strongly addressed. The formal, auditable approval workflow within Git and the strict access controls

on the repository provide powerful evidence to auditors that changes are properly authorized, tested, and documented, and that segregation of duties is enforced.¹¹

- **Health Insurance Portability and Accountability Act (HIPAA):** The HIPAA Security Rule mandates a set of Technical Safeguards to protect electronic Protected Health Information (ePHI). A secure configuration management framework is essential for meeting these requirements:
 - **Access Control (§ 164.312(a)(1)):** This is a required specification. The use of unique user IDs in Git and the CI/CD system, coupled with RBAC/ABAC for repository access, directly fulfills this mandate.⁴⁵
 - **Audit Controls (§ 164.312(b)):** This is a required specification. The immutable Git history and the comprehensive logs from the CI/CD pipeline and target systems provide the mechanism to "record and examine activity" involving systems that process ePHI.⁴⁵
 - **Integrity (§ 164.312(c)(1)):** This is a required specification. The use of cryptographic hashing in Git and the tamper-evident logging techniques discussed in Part III provide strong protection against the improper alteration or destruction of configurations that protect ePHI.⁴⁵
- **Financial Industry Regulatory Authority (FINRA):** FINRA places a heavy emphasis on cybersecurity for its member firms. Its Cybersecurity Checklist, which is derived from the NIST CSF, explicitly requires firms to establish and maintain programs to protect assets and detect breaches.⁴⁷ FINRA audits specifically assess firms on their controls for **access management, system change management, and technology governance.**⁴⁷ A well-documented and enforced framework for agent configuration, particularly one built on GitOps principles, provides a clear and defensible response to these points of regulatory scrutiny.⁴⁹

Ultimately, the evolution of regulatory frameworks, often perceived as lagging behind technology, is implicitly driving the adoption of these more advanced, cryptographically verifiable controls. As the sophistication of digital threats grows, the definition of what constitutes a "sufficient" control for ensuring integrity and auditability is rising. The advanced logging techniques, such as Merkle trees and blockchain anchoring, are becoming the new gold standard for *proving* compliance with the long-standing principles embedded in these regulations. A forward-looking security leader will adopt these modern techniques not just for their superior security properties, but as a proactive measure to meet and exceed the rising bar of regulatory expectation.

Conclusions

The secure management of agent configurations is a multi-faceted challenge that requires a holistic framework built on a foundation of sound principles, robust procedures, and modern technology. A successful strategy cannot be achieved by implementing a single tool or policy in isolation. Instead, it demands the integration of access control, change management, and audit logging into a cohesive system that is both defensible to auditors and resilient against threats.

The analysis yields several key conclusions for leaders tasked with architecting these controls:

1. **The Principle of Least Privilege is Non-Negotiable:** PoLP must be the philosophical starting point for all access decisions. Its rigorous implementation, through a combination of audits, default-deny policies, and continuous monitoring, is the single most effective measure for reducing the attack surface and containing the impact of a potential breach. The choice between RBAC and ABAC to implement PoLP is a strategic one, with significant downstream consequences for administrative overhead, flexibility, and the complexity of the required audit infrastructure.
2. **Multi-Person Approval for Critical Changes is a Mandate:** For any configuration change that carries significant risk, the "maker-checker" or "four-eyes" principle, which enforces Segregation of Duties, is an essential control. This procedural safeguard mitigates both internal fraud and human error. In modern environments, this principle is best implemented not as a manual workflow but as an automated, programmatic guardrail within the development toolchain, such as mandatory, multi-person pull request approvals in a GitOps workflow.
3. **The Audit Trail Must Be Immutable and Comprehensive:** An effective audit trail is more than a log file; it is a complete, context-rich, and trustworthy record of an event's entire lifecycle. This requires not only capturing detailed information about every change but also ensuring the log's integrity through advanced cryptographic techniques like hash chaining, Merkle trees, and digital signatures. For the highest level of assurance, organizations should pursue true immutability by leveraging trusted third-party logging services and anchoring log history in a public ledger.

4. **Modernization is a Security Imperative:** The adoption of IaC and GitOps practices fundamentally transforms configuration management for the better. It shifts controls from slow, manual processes to fast, automated, and inherently auditable workflows. This paradigm allows security to be embedded directly into the development lifecycle, enabling "secure velocity" rather than acting as a bottleneck. The role of the security team evolves from that of a manual gatekeeper to an expert architect of the automated guardrails that make this speed possible.

Ultimately, the goal is to create a system where the secure path is the easiest path. By building a framework that combines the strategic clarity of governance models like NIST and COBIT, the procedural rigor of SoD, and the technical power of modern automation and cryptography, organizations can achieve a state of continuous compliance and proactive security. This allows them to manage their critical agent configurations with confidence, ensuring that these powerful tools remain assets for the business rather than liabilities waiting to be exploited.

Ouvrages cités

1. What is Principle of Least Privilege (PoLP)? | Cato Networks, dernier accès : août 12, 2025, <https://www.catonetworks.com/glossary/principle-of-least-privilege/>
2. What is Principle of Least Privilege (POLP)? - CrowdStrike, dernier accès : août 12, 2025, <https://www.crowdstrike.com/en-us/cybersecurity-101/identity-protection/principle-of-least-privilege-polp/>
3. The Least Privilege Policy Explained - Delinea, dernier accès : août 12, 2025, <https://delinea.com/what-is/least-privilege>
4. What is the Principle of Least Privilege? | NordLayer Learn, dernier accès : août 12, 2025, <https://nordlayer.com/learn/access-control/principle-of-least-privilege/>
5. What Is the Principle of Least Privilege? - Palo Alto Networks, dernier accès : août 12, 2025, <https://www.paloaltonetworks.com/cyberpedia/what-is-the-principle-of-least-privilege>
6. Understanding RBAC, PBAC, and ABAC: Access Control Explained - Ping Identity, dernier accès : août 12, 2025, <https://www.pingidentity.com/en/resources/identity-fundamentals/authorization/authorization-methods.html>
7. ABAC vs. RBAC: What's the difference? - Citrix Blogs, dernier accès : août 12, 2025, <https://www.citrix.com/blogs/2022/05/17/abac-vs-rbac-comparison/>
8. RBAC vs. ABAC: Role-Based & Attribute-Based Access Control Compared | Splunk, dernier accès : août 12, 2025, https://www.splunk.com/en_us/blog/learn/rbac-vs-abac.html
9. Protect | NIST, dernier accès : août 12, 2025,

- <https://www.nist.gov/cyberframework/protect>
10. ISO 27001 Access Control Policy Beginner's Guide (& template) - High Table, dernier accès : août 12, 2025, <https://hightable.io/iso-27001-access-control-policy-ultimate-guide/>
 11. IT General Controls (ITGC): Everything You Need to Know - Scytale, dernier accès : août 12, 2025, <https://scytale.ai/resources/it-general-controls-itgc-everything-you-need-to-know/>
 12. The 5-Step Guide to IT General Controls for SOX Compliance - Scytale, dernier accès : août 12, 2025, <https://scytale.ai/resources/the-5-step-guide-to-it-general-controls-for-sox-compliance/>
 13. What is the four-eyes principle? | UNIDO, dernier accès : août 12, 2025, <https://www.unido.org/overview/member-states/change-management/faq/what-four-eyes-principle>
 14. Maker-checker - Wikipedia, dernier accès : août 12, 2025, <https://en.wikipedia.org/wiki/Maker-checker>
 15. Maker Checker for Better Control on Expense Approvals - Volopay, dernier accès : août 12, 2025, <https://www.volopay.com/expense-management/maker-checker-workflow-for-expense-approvals/>
 16. The Maker-Checker Process And Its Benefits For Your Business - Wallex, dernier accès : août 12, 2025, <https://www.wallex.asia/en-sg/resources/articles/the-maker-checker-process>
 17. Workflow Approval Process: How to Create and Build - SolveXia, dernier accès : août 12, 2025, <https://www.solvexia.com/blog/workflow-approval-process-how-to-create-and-build>
 18. Approval Workflows: Time to Ditch Your Email Ping-Pongs! - Cflow, dernier accès : août 12, 2025, <https://www.cflowapps.com/approval-workflow/>
 19. What is GitOps? - GitLab, dernier accès : août 12, 2025, <https://about.gitlab.com/topics/gitops/>
 20. CI/CD Pipeline Security Best Practices: The Ultimate Guide - Wiz, dernier accès : août 12, 2025, <https://www.wiz.io/academy/ci-cd-security-best-practices>
 21. Change Management Process Checklist using the COBIT Control ..., dernier accès : août 12, 2025, <https://www.dbta.com/DBTA-Downloads/WhitePapers/Change-Management-Process-Checklist-using-the-COBIT-Control-Objectives-6190.pdf>
 22. CI/CD Best Practices - Top 11 Tips for Successful Pipelines - Spacelift, dernier accès : août 12, 2025, <https://spacelift.io/blog/ci-cd-best-practices>
 23. What key elements should an effective audit trail include? | Simple ..., dernier accès : août 12, 2025, <https://sbnsoftware.com/blog/what-key-elements-should-an-effective-audit-trail-include/>
 24. What is Log Tampering? - GeeksforGeeks, dernier accès : août 12, 2025,

- <https://www.geeksforgeeks.org/ethical-hacking/what-is-log-tampering/>
25. A Tamperproof Logging Implementation - Pangea.cloud, dernier accès : août 12, 2025, <https://pangea.cloud/blog/a-tamperproof-logging-implementation/>
 26. Tamper-proof logs | Identification for Development - World Bank ID4D, dernier accès : août 12, 2025, <https://id4d.worldbank.org/guide/tamper-proof-logs>
 27. Efficient Data Structures For Tamper-Evident Logging. - ResearchGate, dernier accès : août 12, 2025, https://www.researchgate.net/publication/221260542_Efficient_Data_Structures_For_Tamper-Evident_Logging
 28. Tamper-Evident Logs - Hacker News, dernier accès : août 12, 2025, <https://news.ycombinator.com/item?id=25995034>
 29. Cybersecurity Framework | NIST, dernier accès : août 12, 2025, <https://www.nist.gov/cyberframework>
 30. Understanding the NIST cybersecurity framework - Federal Trade Commission, dernier accès : août 12, 2025, <https://www.ftc.gov/business-guidance/small-businesses/cybersecurity/nist-framework>
 31. NIST Cybersecurity Framework (CSF) Core Explained - CyberSaint, dernier accès : août 12, 2025, <https://www.cybersaint.io/blog/nist-cybersecurity-framework-core-explained>
 32. The NIST Cybersecurity Framework (CSF) 2.0, dernier accès : août 12, 2025, <https://nvlpubs.nist.gov/nistpubs/CSWP/NIST.CSWP.29.pdf>
 33. What is COBIT Framework ? A Detailed Guide - Sprinto, dernier accès : août 12, 2025, <https://sprinto.com/blog/cobit-framework/>
 34. What is COBIT? Principles and Enablers Explained - AuditBoard, dernier accès : août 12, 2025, <https://auditboard.com/blog/cobit>
 35. Understanding the COBIT Framework: A Comprehensive Guide - The LastPass Blog, dernier accès : août 12, 2025, <https://blog.lastpass.com/posts/cobit-framework>
 36. What is ISO/IEC 27001, The Information Security Standard - ISMS.online, dernier accès : août 12, 2025, <https://www.isms.online/iso-27001/>
 37. Understanding ISO/IEC 27001: Your path to compliance. - Protecht, dernier accès : août 12, 2025, <https://www.protechtgroup.com/en-us/blog/iso-iec-27001-compliance-guide>
 38. ISO 27001:2022 - A.5 Organisational Controls (Access Management) - URM Consulting, dernier accès : août 12, 2025, <https://www.urmconsulting.com/blog/iso-27001-2022-a-5-organisational-controls-access-management>
 39. ISO 27001 Controls Explained: A Guide to Annex A (Updated 2024) | Secureframe, dernier accès : août 12, 2025, <https://secureframe.com/hub/iso-27001/controls>
 40. ISO 27001 Controls Cheat Sheet - Wiz, dernier accès : août 12, 2025, <https://www.wiz.io/academy/iso-27001-controls>
 41. ISO 27001 Change Management: Annex A 8.32 - High Table, dernier accès : août 12, 2025, <https://hightable.io/iso27001-annex-a-8-32-change-management/>
 42. What Is Infrastructure as Code (IaC)? - Zscaler, Inc., dernier accès : août 12, 2025,

<https://www.zscaler.com/resources/security-terms-glossary/what-is-infrastructure-as-code-security>

43. Infrastructure as Code : Best Practices, Benefits & Examples - Spacelift, dernier accès : août 12, 2025, <https://spacelift.io/blog/infrastructure-as-code>
44. GitOps: Managing Infrastructure and Applications with Git - Keitaro, dernier accès : août 12, 2025, <https://www.keitaro.com/insights/2024/10/18/gitops-managing-infrastructure-and-applications-with-git/>
45. HIPAA Security Technical Safeguards - ASHA, dernier accès : août 12, 2025, <https://www.asha.org/practice/reimbursement/hipaa/technicalsafeguards/>
46. What are Technical Safeguards of HIPAA's Security Rule?, dernier accès : août 12, 2025, <https://www.hipaaexams.com/blog/technical-safeguards-security-rule>
47. FINRA - Centraleyes, dernier accès : août 12, 2025, <https://www.centraleyes.com/finra-cybersecurity-checklist/>
48. Cybersecurity Checklist, dernier accès : août 12, 2025, <https://business.cch.com/srd/cybersecurity-checklist.pdf>
49. FINRA Cybersecurity Checklist - RSI Security, dernier accès : août 12, 2025, <https://www.rsisecurity.com/wp-content/uploads/2022/07/FINRA.pdf>